



The perfect agent is coming. We want it too.

Until then — let's ship 5–10× today, and iterate.

We all want the same thing

- A workflow where AI ships features end-to-end. **That goal is right.**
- The fully autonomous version — true 24/7, zero-HITL — is still maturing.
- The opportunity: **5–10× today** with a standardized workflow, and we phase in more autonomy as the tooling catches up.
- **Same destination. Real deliverable now. Earlier wins.**

The path forward, today

Standardize on Valkyrie.

- Design-first workflow, enforced by a hook.
- AFK loops with budget caps for heads-down work — without breaking our company's budget.
- **5–10x more efficient. Today. Same workflow for everyone.**

How it works

```
DESIGN → PRD → ISSUES → TDD  
grill   read  slice   red-green
```

Once per feature, say **"let's build X"** (or type `/valk`). The orchestrator routes you through all four stages — every follow-up prompt continues the active stage. A statusline shows where you are.

Try to skip → the AI refuses.

(No sub-commands to memorize — the orchestrator handles them.)

What's "the hook"?

- A tiny script Claude Code runs **before every prompt you submit**.
- It reads what you typed. If it looks like build/fix/refactor language → it injects a non-negotiable directive: *"use the Valkyrie workflow."*
- Claude can't ignore it. No flag to memorize, no opt-in checkbox.
- **Installed once. Applies to every project, every prompt, every engineer.**

Stage 1 — DESIGN

- AI asks one question at a time, **with its recommended answer.**
- You react: accept, push back, or "you decide."
- ~15 minutes. Front-loads the thinking that wastes hours later.
- **Output:** shared understanding of what to build.

Stage 2 — PRD

- AI synthesizes the grill into a Markdown PRD.
- Saved to `docs/prd/<slug>.md` — reviewable, shareable, version-controlled.
- ~5 minutes for you to read. Edit anything wrong.
- **Output:** the spec. Code review now references it instead of the diff.

Stage 3 — ISSUES

- AI breaks the PRD into **vertical slices** — each cuts through every layer (schema → API → UI → tests).
- Marked **AFK** (autonomous-safe) or **HITL** (needs a human).
- Dependency-ordered so the loop knows what's unblocked.
- **Output:** `issues/0001-*.md` , `0002-*.md` , ... each independently shippable.

Stage 4 — TDD

- Red-green-refactor, one issue at a time.
- AI writes a test → you confirm → AI implements → repeat.
- Tests target public behavior, not implementation details.
- **Output:** working code, tested, traceable to the PRD.

Optional: AFK mode

Walk away after stages 1–3. The loop runs Stage 4 autonomously.

- **Prerequisites:** PRD written, issues queued, working tree clean.
- Run from your shell:

```
afk 10                                # defaults: 4h, $50  
afk 50 --max-hours 12 --max-cost-usd 200 # overnight
```

- Spawns a fresh CLI session per issue. No context drift.
- **Pre-flight gates refuse to launch** unless prerequisites are met.

Day 1

```
git clone https://github.com/BunnyDA0/Valkyrie.git ~/valkyrie && cd $_ && ./install.sh
```

Then in any project, just say:

"Let's add a billing dashboard."

The ask

Make this the team standard.

- Pilot: 3–5 volunteers, week 1.
- Default: PRD-backed PRs by quarter end.
- Measure cycle time. Adjust or scale.